Govagn - One-Page App Summary

What It Is

Govagn is a self-hosted control plane for enterprise AI runtime governance, observability, and control.
It unifies policy enforcement, cost attribution, prompt lifecycle, and runtime visibility for LLM and agent traffic.

Who It Is For

Primary persona: enterprise AI platform engineer or governance owner operating internal LLM workloads in controlled environments.

What It Does

- Collects OTLP telemetry and runtime metadata for LLM and agent executions.
- Provides trace, span, lineage, redaction, and live runtime visibility in the portal.
- Enforces policy and guardrails, including preview and simulation workflows.
- Calculates per-request and per-span costs with pricing rules and budget workflows.
- Manages prompt versioning, promotion, and trace-to-release linkage.
- Supports provider adapters or routing for OpenAI, Anthropic, Google, Vertex AI, and Bedrock.
- Ships self-hosted deployment paths via Docker Compose and Kubernetes or Helm.

How It Works (Repo-Evidenced Architecture)

- Components: agent-sdk (Python instrumentation), collector (OTLP ingest and enrichment), api-gateway (policy, pricing, prompt, eval, budget, audit control plane), portal (React operations UI), PostgreSQL (system of record), Redis (cache and coordination).
- Data flow: app or agent to SDK or OTLP to collector to api-gateway to PostgreSQL and Redis; portal queries api-gateway for operations and admin workflows.
- Local stack also includes Prometheus, Alertmanager, Grafana, and Jaeger for monitoring and diagnostics.

How To Run (Minimal Getting Started)

- Install Docker and Docker Compose.
- From repo root, run: make dev
- Open portal at http://localhost:3000 (API gateway at http://localhost:8080).
- Register at least one provider key, then send instrumented or proxied traffic.
- Initial admin credentials: Not found in repo.

Repo evidence: README.md, docker-compose.yml, Makefile, docs/SETUP_AND_ONBOARDING.md